

MSDS 422 Week 9 Assignment 9: Autoencoder

Build an autoencoder for the Kaggle MNIST training dataset.

- (1) Import and preprocess data
- (2) Set parameters and define autoencoder backpropagation function
- (3) Run a process to train autoencoder
- (4) Visualise what the autoencoder has learned

Take the output of that autoencoder and use it to classify the MNIST test set observations. Submit your results to Kaggle.com.

(1) Import and preprocess data

```
In [39]: import keras
from matplotlib import pyplot as plt
import numpy as np
import pandas as pd
import gzip
%matplotlib inline
from keras.models import Model
from keras.optimizers import RMSprop
from keras.layers import Input,Dense,Flatten,Dropout,merge,Reshape,Conv2D,MaxPooling2D,UpSampling2D,Conv2DTranspose
from keras.layers.normalization import BatchNormalization
from keras.models import Model,Sequential
from keras.callbacks import ModelCheckpoint
from keras.optimizers import Adadelta, RMSprop,SGD,Adam
from keras import regularizers
from keras import backend as K
from keras.utils import to_categorical
import tensorflow as tf
```

```
In [40]: df= pd.read_csv("Mtrain.csv")
df_test= pd.read_csv("Mtest.csv")
```

```
In [41]: y = df['label'].values
X = df.drop(['label'], axis=1).values.reshape(-1, 28, 28,1)/255.

train_Y_one_hot = keras.utils.to_categorical(y)

X_test = df_test.values.reshape(-1, 28, 28,1)/255.

from sklearn.model_selection import train_test_split
train_X,valid_X,train_ground,valid_ground = train_test_split(X, X, test_size=0.2, random_state=13)

batch_size = 64
epochs = 1
inChannel = 1
x, y = 28, 28
input_img = tf.keras.Input(shape = (x, y, inChannel))
num_classes = 10
```

(2) Set parameters and define autoencoder backpropagation function

```
In [42]: def encoder(input_img):
conv1 = keras.layers.Conv2D(32, (3, 3), activation='relu', padding='same')(input_img) #28 x 28 x 32
conv1 = keras.layers.BatchNormalization()(conv1)
conv1 = keras.layers.Conv2D(32, (3, 3), activation='relu', padding='same')(conv1)
conv1 = keras.layers.BatchNormalization()(conv1)
pool1 = keras.layers.MaxPool2D(pool_size=(2, 2))(conv1) #14 x 14 x 32
conv2 = keras.layers.Conv2D(64, (3, 3), activation='relu', padding='same')(pool1) #14 x 14 x 64
conv2 = keras.layers.BatchNormalization()(conv2)
conv2 = keras.layers.Conv2D(64, (3, 3), activation='relu', padding='same')(conv2)
conv2 = keras.layers.BatchNormalization()(conv2)
pool2 = keras.layers.MaxPool2D(pool_size=(2, 2))(conv2) #7 x 7 x 64
conv3 = keras.layers.Conv2D(128, (3, 3), activation='relu', padding='same')(pool2) #7 x 7 x 128 (small and thick)
conv3 = keras.layers.BatchNormalization()(conv3)
conv3 = keras.layers.Conv2D(128, (3, 3), activation='relu', padding='same')(conv3)
conv3 = keras.layers.BatchNormalization()(conv3)
conv4 = keras.layers.Conv2D(256, (3, 3), activation='relu', padding='same')(conv3) #7 x 7 x 256 (small and thick)
conv4 = keras.layers.BatchNormalization()(conv4)
conv4 = keras.layers.Conv2D(256, (3, 3), activation='relu', padding='same')(conv4)
conv4 = keras.layers.BatchNormalization()(conv4)
return conv4

def decoder(conv4):
conv5 = keras.layers.Conv2D(128, (3, 3), activation='relu', padding='same')(conv4) #7 x 7 x 128
conv5 = keras.layers.BatchNormalization()(conv5)
conv5 = keras.layers.Conv2D(128, (3, 3), activation='relu', padding='same')(conv5)
conv5 = keras.layers.BatchNormalization()(conv5)
conv6 = keras.layers.Conv2D(64, (3, 3), activation='relu', padding='same')(conv5) #7 x 7 x 64
conv6 = keras.layers.BatchNormalization()(conv6)
conv6 = keras.layers.Conv2D(64, (3, 3), activation='relu', padding='same')(conv6)
conv6 = keras.layers.BatchNormalization()(conv6)
up1 = keras.layers.UpSampling2D((2,2))(conv6) #14 x 14 x 64
conv7 = keras.layers.Conv2D(32, (3, 3), activation='relu', padding='same')(up1) # 14 x 14 x 32
conv7 = keras.layers.BatchNormalization()(conv7)
conv7 = keras.layers.Conv2D(32, (3, 3), activation='relu', padding='same')(conv7)
conv7 = keras.layers.BatchNormalization()(conv7)
up2 = keras.layers.UpSampling2D((2,2))(conv7) # 28 x 28 x 32
decoded = keras.layers.Conv2D(1, (3, 3), activation='sigmoid', padding='same')(up2) # 28 x 28 x 1
return decoded
```

(3) Run a process to train autoencoder

```
In [43]: autoencoder = keras.Model(input_img, decoder(encoder(input_img)))
autoencoder.compile(loss='mean_squared_error', optimizer = keras.optimizers.RMSprop())

#Train the model
autoencoder_train = autoencoder.fit(train_X, train_ground, batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(valid_X, valid_ground))

#Take a look at the Loss
loss = autoencoder_train.history['loss']
val_loss = autoencoder_train.history['val_loss']
epochs = range(1)
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()

#Save the model
autoencoder.save_weights('autoencoder.h5')

# Predict with the encoder weights you just trained
# Change the labels from categorical to one-hot encoding
##train_Y_one_hot = keras.utils.to_categorical(train_labels)
##test_Y_one_hot = keras.utils.to_categorical(test_labels)

# Display the change for category label using one-hot encoding
##print('Original label:', train_labels[0])
##print('After conversion to one-hot:', train_Y_one_hot[0])

train_X,valid_X,train_label,valid_label = train_test_split(X,train_Y_one_hot,test_size=0.2,random_state=13)

#Check the shape
train_X.shape,valid_X.shape,train_label.shape,valid_label.shape


#fully connected layers
def fc(enco):
    flat = tf.keras.layers.Flatten()(enco)
    den = tf.keras.layers.Dense(128, activation='relu')(flat)
    out = tf.keras.layers.Dense(num_classes, activation='softmax')(den)
    return out

encode = encoder(input_img)
full_model = keras.Model(input_img,fc(encode))
```

Model: "functional_15"

Layer (type)	Output Shape	Param #
=====		
input_10 (InputLayer)	[(None, 28, 28, 1)]	0
conv2d_64 (Conv2D)	(None, 28, 28, 32)	320
batch_normalization_60 (Batch Normalization)	(None, 28, 28, 32)	128
conv2d_65 (Conv2D)	(None, 28, 28, 32)	9248
batch_normalization_61 (Batch Normalization)	(None, 28, 28, 32)	128
max_pooling2d_13 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_66 (Conv2D)	(None, 14, 14, 64)	18496
batch_normalization_62 (Batch Normalization)	(None, 14, 14, 64)	256
conv2d_67 (Conv2D)	(None, 14, 14, 64)	36928
batch_normalization_63 (Batch Normalization)	(None, 14, 14, 64)	256
max_pooling2d_14 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_68 (Conv2D)	(None, 7, 7, 128)	73856
batch_normalization_64 (Batch Normalization)	(None, 7, 7, 128)	512
conv2d_69 (Conv2D)	(None, 7, 7, 128)	147584
batch_normalization_65 (Batch Normalization)	(None, 7, 7, 128)	512
conv2d_70 (Conv2D)	(None, 7, 7, 256)	295168
batch_normalization_66 (Batch Normalization)	(None, 7, 7, 256)	1024
conv2d_71 (Conv2D)	(None, 7, 7, 256)	590080
batch_normalization_67 (Batch Normalization)	(None, 7, 7, 256)	1024
conv2d_72 (Conv2D)	(None, 7, 7, 128)	295040
batch_normalization_68 (Batch Normalization)	(None, 7, 7, 128)	512
conv2d_73 (Conv2D)	(None, 7, 7, 128)	147584
batch_normalization_69 (Batch Normalization)	(None, 7, 7, 128)	512

conv2d_74 (Conv2D)	(None, 7, 7, 64)	73792
batch_normalization_70 (Batch Normalization)	(None, 7, 7, 64)	256
conv2d_75 (Conv2D)	(None, 7, 7, 64)	36928
batch_normalization_71 (Batch Normalization)	(None, 7, 7, 64)	256
up_sampling2d_11 (UpSampling2D)	(None, 14, 14, 64)	0
conv2d_76 (Conv2D)	(None, 14, 14, 32)	18464
batch_normalization_72 (Batch Normalization)	(None, 14, 14, 32)	128
conv2d_77 (Conv2D)	(None, 14, 14, 32)	9248
batch_normalization_73 (Batch Normalization)	(None, 14, 14, 32)	128
up_sampling2d_12 (UpSampling2D)	(None, 28, 28, 32)	0
conv2d_78 (Conv2D)	(None, 28, 28, 1)	289
=====		
Total params: 1,758,657		
Trainable params: 1,755,841		
Non-trainable params: 2,816		
=====		
525/525 [=====] - 350s 666ms/step - loss: 0.0237 - val_loss: 0.0133		




(4) Visualise what the autoencoder has learned

```
In [67]: #Visualize the reconstructions
def plot_image(image):
    plt.imshow(image,cmap='binary')
plt.axis('off')

def show_reconstructions(model,n_images=5):
    reconstructions = model.predict(valid_X[:n_images])
    fig = plt.figure(figsize=(n_images*1.5,3))
    for image_index in range(n_images):
        plt.subplot(2,n_images,1 + image_index)
        plot_image(valid_X[image_index].reshape(28,28))
        plt.subplot(2,n_images,1+ n_images + image_index)
        plot_image(reconstructions[image_index].reshape(28,28))

show_reconstructions(autoencoder)
```



MNIST Test

```
In [45]: full_model.compile(loss=keras.losses.categorical_crossentropy, optimizer=keras.optimizers.Adam(),metrics=['accuracy'])

print(full_model.summary())

#Train the model
classify_train = full_model.fit(train_X, train_label, batch_size=64,epochs=1,verbose=1,validation_data=(valid_X, valid_label))

#Save the model
full_model.save_weights('autoencoder_classification.h5')

#plot accuracy
accuracy = classify_train.history['accuracy']
val_accuracy = classify_train.history['val_accuracy']
loss = classify_train.history['loss']
val_loss = classify_train.history['val_loss']
epochs = range(len(accuracy))
plt.plot(epochs, accuracy, 'bo', label='Training accuracy')
plt.plot(epochs, val_accuracy, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()

#Predictions
predicted_classes = full_model.predict(X_test)
predicted_classes = np.argmax(np.round(predicted_classes),axis=1)
```


Model: "functional_17"

Layer (type)	Output Shape	Param #
=====		
input_10 (InputLayer)	[(None, 28, 28, 1)]	0
conv2d_79 (Conv2D)	(None, 28, 28, 32)	320
batch_normalization_74 (Batch Normalization)	(None, 28, 28, 32)	128
conv2d_80 (Conv2D)	(None, 28, 28, 32)	9248
batch_normalization_75 (Batch Normalization)	(None, 28, 28, 32)	128
max_pooling2d_15 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_81 (Conv2D)	(None, 14, 14, 64)	18496
batch_normalization_76 (Batch Normalization)	(None, 14, 14, 64)	256
conv2d_82 (Conv2D)	(None, 14, 14, 64)	36928
batch_normalization_77 (Batch Normalization)	(None, 14, 14, 64)	256
max_pooling2d_16 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_83 (Conv2D)	(None, 7, 7, 128)	73856
batch_normalization_78 (Batch Normalization)	(None, 7, 7, 128)	512
conv2d_84 (Conv2D)	(None, 7, 7, 128)	147584
batch_normalization_79 (Batch Normalization)	(None, 7, 7, 128)	512
conv2d_85 (Conv2D)	(None, 7, 7, 256)	295168
batch_normalization_80 (Batch Normalization)	(None, 7, 7, 256)	1024
conv2d_86 (Conv2D)	(None, 7, 7, 256)	590080
batch_normalization_81 (Batch Normalization)	(None, 7, 7, 256)	1024
flatten_4 (Flatten)	(None, 12544)	0
dense_8 (Dense)	(None, 128)	1605760
dense_9 (Dense)	(None, 10)	1290
=====		

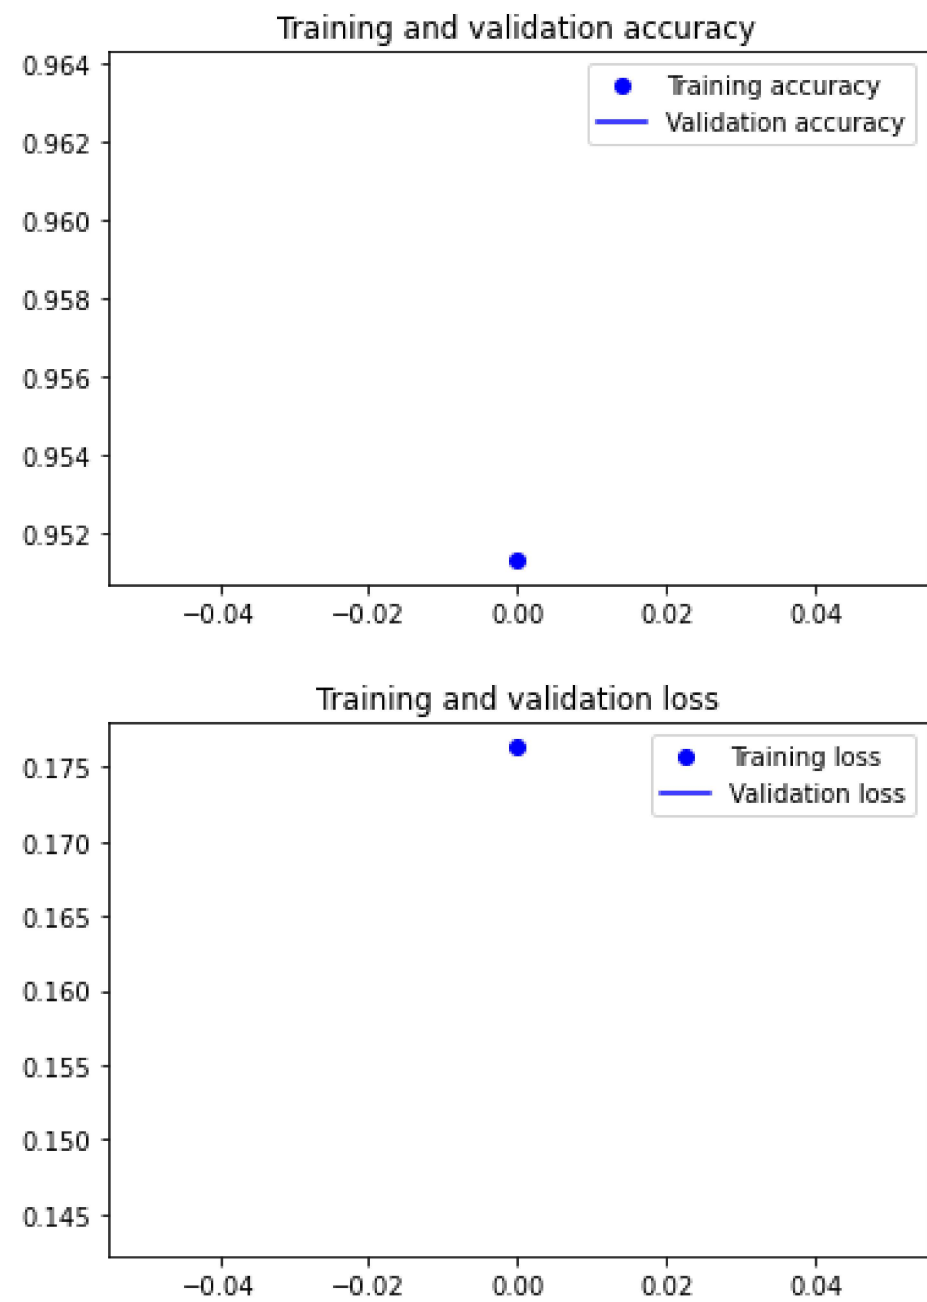
Total params: 2,782,570

Trainable params: 2,780,650

Non-trainable params: 1,920

None

525/525 [=====] - 240s 456ms/step - loss: 0.1763 - accuracy: 0.9513 - val_loss: 0.1438 - val_accuracy: 0.9637



	ImageId	Label
0	1	2
1	2	0
2	3	9
3	4	0
4	5	3
...
27995	27996	9
27996	27997	7
27997	27998	3
27998	27999	9
27999	28000	2

[28000 rows x 2 columns]

PermissionError Traceback (most recent call last)

<ipython-input-45-e8886bb2da73> in <module>

35

36 print(submission)

---> 37 submission.to_csv("submission.csv",index=False)

~\Anaconda3\lib\site-packages\pandas\core\generic.py in to_csv(self, path_or_buf, sep, na_rep, float_format, columns, header, index, index_label, mode, encoding, compression, quoting, quotechar, line_terminator, chunksize, date_format, doublequote, escapechar, decimal)

3202 decimal=decimal,

3203)

-> 3204 formatter.save()

3205

3206 if path_or_buf is None:

~\Anaconda3\lib\site-packages\pandas\io\formats\csvs.py in save(self)

182 close = False

183 else:

--> 184 f, handles = get_handle(

185 self.path_or_buf,

186 self.mode,

~\Anaconda3\lib\site-packages\pandas\io\common.py in get_handle(path_or_buf, mode, encoding, compression, memory_map, is_text)

426 if encoding:

427 # Encoding

--> 428 f = open(path_or_buf, mode, encoding=encoding, newline="")

429 elif is_text:

430 # No explicit encoding

PermissionError: [Errno 13] Permission denied: 'submission.csv'

Kaggle Score

```
In [68]: results = pd.Series(predicted_classes,name="Label")
submission = pd.concat([pd.Series(range(1,28001),name = "ImageId"),results],axis = 1)

print(submission)
submission.to_csv("submission.csv",index=False)
```

	ImageId	Label
0	1	2
1	2	0
2	3	9
3	4	0
4	5	3
...
27995	27996	9
27996	27997	7
27997	27998	3
27998	27999	9
27999	28000	2

[28000 rows x 2 columns]